

Wenxin Ding

Email: wenxind@uchicago.edu

Website: wenxind.github.io

RESEARCH INTEREST

My research interest lies in machine learning security and privacy. Specifically, I focus on bridging the gap between theoretical understanding and empirical practice. Recently, I have been working on problems regarding vulnerabilities of diffusion models.

EDUCATION

University of Chicago

Ph.D. in Computer Science

Advisors: Prof. Heather Zheng and Prof. Ben Y. Zhao

Chicago, IL

Sep 2021 – June 2026

Carnegie Mellon University

M.S. in Computer Science – Research Thesis

Advisors: Prof. Nihar Shah and Prof. Weina Wang

Pittsburgh, PA

Aug 2021

B.S. in Computer Science and B.S. in Mathematical Sciences

Minor in Computational Finance

May 2020

PUBLICATIONS

Conference

- **Wenxin Ding**, Cathy Li, Shawn Shan, Ben Y. Zhao, Haitao Zheng. “Understanding Implosion in Text-to-Image Generative Models.” *2024 ACM SIGSAC Conference on Computer and Communications Security (CCS)*.
- Shawn Shan, **Wenxin Ding**, Josephine Passananti, Haitao Zheng, Ben Y. Zhao. “Prompt-Specific Poisoning Attacks on Text-to-Image Generative Models.” *2024 IEEE Symposium on Security and Privacy (SP)*.
- **Wenxin Ding**, Arjun Nitin Bhagoji, Ben Y. Zhao, and Haitao Zheng. “Towards Scalable and Robust Model Versioning.” *2nd IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*.
- Sihui Dai*, **Wenxin Ding***, Arjun Nitin Bhagoji, Daniel Cullina, Ben Y. Zhao, Haitao Zheng, and Prateek Mittal. “Characterizing the Optimal 0-1 Loss for Multi-class Classification with a Test-time Attacker.” *2023 Advances in Neural Information Processing Systems (NeurIPS)*.
Spotlight
- Shawn Shan, **Wenxin Ding**, Emily Wenger, Haitao Zheng, and Ben Y. Zhao. “Post-breach recovery: Protection against white-box adversarial examples for leaked DNN models.” *2022 ACM SIGSAC Conference on Computer and Communications Security (CCS)*.

- **Wenxin Ding**, Gautam Kamath, Weina Wang, and Nihar B. Shah. “Calibration with privacy in peer review.” *2022 IEEE International Symposium on Information Theory (ISIT)*.

Workshop

- **Wenxin Ding**, Nihar B. Shah, and Weina Wang. “On the privacy-utility tradeoff in peer-review data analysis.” *2021 AAAI Privacy-Preserving Artificial Intelligence (PPAI) workshop*.
Spotlight

TEACHING EXPERIENCE

Teaching Assistant

University of Chicago

- CMSC 25300/35300 Mathematical Foundations of Machine Learning

Carnegie Mellon University

- 15110 Principles of Computing (Head Teaching Assistant)
- 15213 Introduction to Computer Systems
- 15440 Distributed Systems

Mentor

Strong Women Strong Girls, Pittsburgh, PA

SERVICE

- Program Committee, 2024 ACM Workshop on Artificial Intelligence and Security (AISec)
- Reviewer, 2024 The Conference on Uncertainty in Artificial Intelligence (UAI)
- Volunteer, 2022 ACM Conference on Computer and Communications Security (CCS)

AWARDS

- 2021 University of Chicago Eckhardt Scholar
- 2020 Carnegie Mellon University Senior Leadership Recognition
- 2019 Mark Stehlik SCS Alumni Undergraduate Impact Scholarship
- 2017 William Lowell Putnam Mathematical Competition (Rank: 255 / 4638)